

hw3

October 3, 2018

1 Meteo 515 – Assignment 3 – Linear Regression

```
In [1]: from __future__ import division, print_function
        #from collections import OrderedDict
        from itertools import chain
        #import datetime as dt

        #import matplotlib.dates as mdates
        import matplotlib.pyplot as plt
        import numpy as np
        import pandas as pd
        import scipy.optimize as so
        import scipy.stats as ss
        #import sklearn as skl
        import statsmodels.api as sm
```

```
In [2]: plt.style.use('seaborn-darkgrid')
        %matplotlib notebook
```

1.1 Load the data

We are using [this dataset](https://esrl.noaa.gov/psd/data/timeseries/AMO/), the not-detrended and unsmoothed AMO index from NOAA ESRL. More info here: <https://esrl.noaa.gov/psd/data/timeseries/AMO/>

Note: Currently (27-Sep-18 10:00 EST) the data on the site is messed up for 1980 onwards, though it was fine when I downloaded it the day before.

```
In [3]: amo_fpath = './data/amon.us.long.mean.data'

amo_raw = np.genfromtxt(amo_fpath, skip_header=1, skip_footer=7)[: ,1:]

with open(amo_fpath, 'r') as f: yr_range = f.readline().split()
t_amo = pd.date_range(start='{}/01/01'.format(yr_range[0]), freq='MS', periods=amo_raw.size)
amo_ndt_us = amo_raw.reshape((amo_raw.size,))

amo = pd.DataFrame({'amo': amo_ndt_us}, index=t_amo)
amo['julian_date'] = amo.index.to_julian_date() # Julian Date (decimal d
amo['t_elapsed'] = amo.julian_date - amo.julian_date.iloc[0] # elapsed time (decimal
amo['year'] = amo.index.year # integer year for each
```

```

amo['decyear'] = amo.index.year + (amo.index.month-1)/12      # decimal year

amo[amo == -99.99] = np.nan
amo.dropna(inplace=True)

#> don't include 2018 in the annual means, since 4 months are missing
grouped = amo.loc[amo.year<=2017, :].groupby(pd.Grouper(freq='A'))
amo_annual_mean = grouped.mean()
amo_annual_mean['sem_annual'] = grouped.amo.std().values / np.sqrt(12) # stdev of the
# ~ note that this gives indices that are the last day of the year
#   and {year}.46 for decimal year, since initially the datetimes are first day of the

In [4]: #> plotting functions!
figsize_lin_reg = (9, 4.0)
figsize_res = (8, 3)
#degC = u'\u00B0C'
degC = u' (\u00B0C)'

def res_plot(y, y_hat, df_res=2, figid='ts_res'):
    """Plot the residuals
    could pass the OLS result in, but not doing that currently...
    """
    f, a = plt.subplots(figsize=figsize_res, num=figid)

    y_bar = y.mean()
    SSE = np.sum((y-y_hat)**2) # sum of squared error
    SSR = np.sum((y_hat-y_bar)**2) # residual sum of squares
    SST = np.sum((y-y_bar)**2)
    assert( np.isclose(SST, SSE+SSR) )
    #MSE = SSE/(res.df_resid) # MSE == sample variance of the residuals
    MSE = SSE/df_res
    #assert( np.isclose(MSE, res.mse_resid) ) # res.mse_model == MSR; F = MSR/MSE
    s = '''
    SSE = {:.3g}
    SSR = {:.3g}
    MSE = {:.4g}
    '''.format(SSE, SSR, MSE)

    a.plot(t_plot, y-y_hat, '.-', c='purple', ms=6, lw=1, label=s)
    #a.set_ylabel('residual AMO index {:.s}'.format(degC))
    a.set_ylabel('residual AMO index' + degC)
    a.text(1.01, 0.5, s,
          va='center', ha='left', transform=a.transAxes)
    #a.legend(loc='center left', bbox_to_anchor=(0.99, 0.5), labelspace=1.5)

    f.tight_layout(rect=(0, 0, 0.85, 1.0))

    return f

```

```

def lin_reg_plot(t_plot, y, y_hat, s_model, figid='ts'):
    """Plot the data and least-squares regression result"""
    f, a = plt.subplots(figsize=figsize_lin_reg, num=figid)

    #xbar, se = y, y.std()
    #a.fill_between(amo_annual_mean.index, xbar-1.*se, xbar+1.*se, alpha=0.3, label='S')
    a.plot(t_plot, y, 'b.-', alpha=0.6, ms=6, lw=1, label='annual means')
    a.plot(t_plot, y_hat, '-', c='r', lw=2, label=s)

    #a.text(0.02, 0.98, s,
    #      va='top', ha='left', transform=a.transAxes)

    a.set_ylabel('AMO index' + degC)
    a.legend(loc='center left', bbox_to_anchor=(0.99, 0.5), labelspace=1.5)

    f.tight_layout()

    return f

```

1.2 1) Simple linear regression with year as predictor

Note that using the monthly data vs the annual means gives slightly different answers, mainly for the error. Change df_reg to see.

```

In [5]: df_reg = amo_annual_mean # {amo_annual_mean, amo}
        t_reg = df_reg.year
        t_plot = df_reg.index

        res = sm.OLS(df_reg['amo'], sm.add_constant(t_reg), ).fit()

        print(res.summary()) # note: can plot res using `sm.graphics.abline_plot(model_result.

        param_lines = '\n'.join([r'$\beta_{:d} = \{:.3g\} \ ( \{:.3g\})$'.format(i) for i in range(1,
        s = r'$y=\beta_1 t + \beta_0$' + '\n\n' + param_lines
        s = s.format(*chain.from_iterable(zip(res.params, res.bse)))

        y = df_reg['amo'].values
        y_hat = t_reg*res.params[1]+res.params[0]

        f1 = lin_reg_plot(t_plot, y, y_hat, s, 'ts_ols1')

        f1b = res_plot(y, y_hat, df_res=res.df_resid, figid='ts_ols1_res');

```

OLS Regression Results

```

=====
Dep. Variable:          amo    R-squared:                0.340

```

```

Model:                OLS      Adj. R-squared:      0.336
Method:               Least Squares    F-statistic:      82.56
Date:                 Wed, 03 Oct 2018    Prob (F-statistic): 3.73e-16
Time:                 21:33:48    Log-Likelihood:    47.402
No. Observations:     162    AIC:                -90.80
Df Residuals:         160    BIC:                -84.63
Df Model:              1
Covariance Type:      nonrobust

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const         15.3469      0.591     25.952      0.000      14.179      16.515
year           0.0028      0.000      9.086      0.000       0.002       0.003
=====
Omnibus:                 3.980    Durbin-Watson:           0.612
Prob(Omnibus):            0.137    Jarque-Bera (JB):        2.915
Skew:                    -0.182    Prob(JB):                0.233
Kurtosis:                 2.454    Cond. No.                8.02e+04
=====

```

Warnings:

```

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
[2] The condition number is large, 8.02e+04. This might indicate that there are
strong multicollinearity or other numerical problems.

```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.Javascript object>
```

```
<IPython.core.display.HTML object>
```

1.2.1 Discussion:

Our R^2 here is pretty low. The value (0.34) indicates that our model explains only 34% of the variation in AMO index. In the plot of the residuals, we see a clear oscillatory pattern, which next we add to our model to see how that improves things...

1.3 2) Adding a prescribed oscillation

```

In [6]: phase = -21 # years
        period = 65 # in years

```

```

y = df_reg['amo'] # AMO index

x0 = np.ones(df_reg['amo'].shape) # intercept
x1 = df_reg['year'].values # year
x2 = np.sin(2*np.pi/period*(x1-phase)) # oscillation with prescribed period and phase

#X = np.vstack((x2, x1, x0)).T
X = pd.DataFrame(data={'0-const': x0, '1-year': x1, '2-osc_amp': x2}, index=df_reg.index)

res = sm.OLS(y, X).fit()

y_hat = np.dot(X, res.params)

print(res.summary())

param_lines = '\n'.join([r'$\beta_{:d} = \{:.3g\} \ (\{:.3g\})$'.format(i) for i in range(3)]
s = r'''
$y = \beta_2 \sin\left(\frac{2 \pi}{\mathrm{period}} (t - \mathrm{phase})\right) + \beta_1 t + \beta_0$

period = {d}
phase = {d}
''' + param_lines
s = s.format(period, phase, *chain.from_iterable(zip(res.params, res.bse)))

f2 = lin_reg_plot(t_plot, y, y_hat, s, 'ts_ols2')

f2b = res_plot(y, y_hat, df_res=res.df_resid, figid='ts_ols2_res');

```

OLS Regression Results

```

=====
Dep. Variable:          amo      R-squared:                0.655
Model:                  OLS      Adj. R-squared:           0.651
Method:                 Least Squares      F-statistic:         151.2
Date:                   Wed, 03 Oct 2018    Prob (F-statistic):    1.65e-37
Time:                   21:33:48           Log-Likelihood:       99.994
No. Observations:       162              AIC:                 -194.0
Df Residuals:           159              BIC:                 -184.7
Df Model:                2
Covariance Type:        nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
0-const	15.4580	0.429	36.044	0.000	14.611	16.305
1-year	0.0027	0.000	12.234	0.000	0.002	0.003
2-osc_amp	0.1779	0.015	12.056	0.000	0.149	0.207

```

=====
Omnibus:                5.696      Durbin-Watson:           1.167

```

Prob(Omnibus):	0.058	Jarque-Bera (JB):	3.405
Skew:	0.153	Prob(JB):	0.182
Kurtosis:	2.359	Cond. No.	8.03e+04

=====

Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 8.03e+04. This might indicate that there are strong multicollinearity or other numerical problems.

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

1.3.1 Discussion:

Adding a prescribed oscillation to our model has greatly improved the fit: our R^2 has increased by about a factor of 2, and MSE correspondingly decreased by about the same. The plot of the residuals is now more like what we would like to see (looking just noisy). However, the values in the first few decades seem to be noticeably and consistently a bit higher than later years, suggesting that maybe the increasing trend didn't start until sometime after 1857. In (3) we will add that to the model and see how things improve...

But first:

1.4 2.5) Find a better oscillation using optimization

```
In [7]: def fit1(t, slope, const, amp, period, phase):
        return amp*np.sin(2*np.pi/period*(t-phase)) + slope*t + const

popt, pcov = so.curve_fit(fit1, df_reg.year, df_reg.amo,
                          p0=(0.0027, 15.5, 0.18, 65, 75))

popt[-1] = popt[-1] ## popt[-2]

y = df_reg.amo
y_hat = fit1(df_reg.year, *popt)

#print(res.summary())

s = r'''
```

```

$y = \beta_2 \sin\left( \frac{2 \pi}{\mathrm{period}} \right) (t - \mathrm{phase}) \sin\left( \frac{2 \pi}{\mathrm{period}} \right) t + \beta_1 t + \beta_0$

$\beta_1 = {:.4g}$
$\beta_0 = {:.4g}$
$\beta_2 = {:.4g}$
period = {:.4g}
phase = {:.4g}
''' .format(*popt)

f2p5 = lin_reg_plot(t_plot, y, y_hat, s, 'ts_ols2.5')

n = df_reg.index.size
f2p5b = res_plot(y, y_hat, df_res=n-5, figid='ts_ols2.5_res');

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```

1.4.1 Discussion:

The improvement in fit is more modest this time, and in fact pretty similar to what we will see below in (3).

1.5 3) Same as (2), but linear trend predictor starts in 1900

Also, we have the option of using the period and phase from the optimized fit in (2.5), but this turned out to not be the best option, since that fit did not include the specification that year only starts as a predictor in 1900, so instead we use the same values as in (2)

Note that the following code for conf and prediction intervals is strictly valid only for **simple linear regression, not multiple**.

```

#> compute confidence and prediction intervals for the predicted
# x: year
# y: AMO
n = len(df_reg.index)
x = df_reg['year'].values
x_bar = x.mean()
s_x = np.sqrt( np.sum((x-x_bar)**2) / (n-1) ) # x.std() does population version!!
y_hat = np.dot(X, fit.params) # should give same as fit.predict(X).values, or fit.fittedvalue
resid = y - y_hat

```